

Neural Information Retrieval

Neural Models

Last lecture

- ❖ You have learned the fundamentals of Deep Learning, with special interest in text processing, where we address the critical problem of text representation.
 - DL -> Neural networks -> Mathematical model -> Learn through optimization
 - Distributed representations
 - Static
 - Contextualized

This lecture

❖ Introduction

- “New” conceptual framework for IR/NIR


❖ Fundamentals in NIR

- The two Neural Architectures
- Retrievers vs Rerankers
- How to train

❖ Literature examples

- Timeline
- Shallow models
- Deep models

Introduction



“New” Conceptual framework for IR

- ❖ Generic formulation that can describe **any** retrieval/ranking model.

$$f(q, D) = \arg \operatorname{top-k}_{d \in D} \phi(\eta_q(q), \eta_d(d))$$

“New” Conceptual framework for IR

- ❖ Generic formulation that can describe **any** retrieval/ranking model.

$$f(q, D) = \arg \operatorname{top-k}_{d \in D} \phi(\eta_q(q), \eta_d(d))$$

❖ With BM25

$\phi(.,.) = \text{inner-product}$

$$\eta_q(q) = \frac{(k_3 + 1) * \vec{q}}{k_3 + \vec{q}},$$

$$\eta_d(d) = IDF(d) * \frac{(k_1 + 1) * \vec{d}}{k_1 * ((1 - b) + b * \frac{|\vec{d}|}{avgdl}) + \vec{d}}$$

Neural Information Retrieval

$$f(q, D) = \arg \max_{d \in D} \phi(\eta_q(q), \eta_d(d))$$

- ❖ Use labelled data to teach an AI model how to search for information!
 - The aim is to let the AI model learn an internal notion of relevance without human intervention



Neural Information Retrieval

Up until now we only used Retrievers


❖ Retrievers:

- Focus on efficiently finding relevant documents from a large corpus
- Optimize for speed and recall

❖ Rerankers:

- Focus on precisely ordering already retrieved documents
- Optimize for precision and ranking quality
- Can analyse deeper semantic relationships between query and documents
- Much more expensive to run when compared to Retrievers
- Work over a query-document pair.


❖ Hybrid approach: Two-stage Retriever




The two architectures for neural IR

$$f(q, D) = \arg \max_{d \in D} \phi(\eta_q(q), \eta_d(d))$$

Representation-Based



Interaction-Based



The two architectures for neural IR

Representation-Based

Ranking

Strengths:

- Efficient Indexing and Retrieval
- Scalability

Weaknesses:

- Poor performance (hard to beat BM25)

Normally, used as **retrieval** models

Interaction-Based

Ranking

Strengths:

- Great performance

Weaknesses:


- Poor efficiency (Impossible to use as a retrieval model)

Normally, used as **reranker** models.

How to train a neural IR model?

❖ Training Dataset

- Collection of queries with a list of **relevant** and **unrelevant** documents.



```
{ "question": "Is Hirschsprung disease a mendelian or a multifactorial disorder?", "goldstandard_documents": [ "PMID:15858239", "PMID:15829955", "PMID:6650562", "PMID:23888072" ] }, { "question": "List signaling molecules (ligands) that interact with the receptor EGFR?", "goldstandard_documents": [ "PMID:21514161", "PMID:23212918", "PMID:23888072" ] }, { "question": "Is the protein Papilin secreted?", "goldstandard_documents": [ "PMID:3320045", "PMID:7515725", "PMID:20805556", "PMID:19297413", "PMID:19724244", "PMID:23888072" ] }, { "question": "Are long non coding RNAs spliced?", "goldstandard_documents": [ "PMID:22955988", "PMID:22955974", "PMID:24285305", "PMID:22707570", "PMID:21622663", "PMID:23888072" ] }, { "question": "Is RANKL secreted from the cells?", "goldstandard_documents": [ "PMID:21618594", "PMID:23698708", "PMID:23827649", "PMID:22901753", "PMID:22948539", "PMID:23888072" ] }, { "question": "Which miRNAs could be used as potential biomarkers for epithelial ovarian cancer?", "goldstandard_documents": [ "PMID:21345725", "PMID:22246341", "PMID:23888072" ] }, { "question": "Which acetylcholinesterase inhibitors are used for treatment of myasthenia gravis?", "goldstandard_documents": [ "PMID:21815707", "PMID:21845054", "PMID:23888072" ] }, { "question": "Has Denosumab (Prolia) been approved by FDA?", "goldstandard_documents": [ "PMID:22826702", "PMID:21208140", "PMID:21113693", "PMID:21129866", "PMID:23888072" ] }, { "question": "List the human genes encoding for the dishevelled proteins?", "goldstandard_documents": [ "PMID:12883684", "PMID:24040443", "PMID:19618470", "PMID:8817300", "PMID:23888072" ] }, { "question": "Name synonym of Acrokeratosis paraneoplastica.", "goldstandard_documents": [ "PMID:3819049", "PMID:18775590", "PMID:6225397", "PMID:22146896", "PMID:7790000", "PMID:23888072" ] }
```

NIR Training – Negative Sampling

- ❖ Finding non-matching document query pairs to contrast with the relevant ones.
 - **Fully random (Random Sampling)**: Randomly samples documents from the collection that are not relevant to the question.
 - **Static Hard (Hard Negatives Sampling)**: Find documents that are similar but irrelevant to the query. For instance, consider the documents retrieved from BM25 as negatives.
 - Positive in BM25 that are not in the gold standard
 - **Dynamic Hard (Curriculum Negative Sampling)**: Start with easy negatives (random) during early training and then start to introduce harder negatives as training progresses.

How to train a neural IR model?

❖ Training Dataset

- Collection of queries with a list of **relevant** and **unrelevant** documents.

Positive example

Negative example

And if we do not have positive documents?

Synthetic data generation with LLMs

```
{
  "question": "Is Hirschsprung disease a mendelian or a multifactorial disorder?",
  "goldstandard_document": "Hirschsprung disease is a multifactorial disorder."
},
{
  "question": "List signaling molecules (ligands) that interact with the receptor EGFR?",
  "goldstandard_documents": [
    "EGFR is a receptor tyrosine kinase that is activated by a variety of ligands, including EGF, TGF-α, and heregulin."
  ]
},
{
  "question": "Is the protein Papilin secreted?",
  "goldstandard_documents": [
    "Papilin is a secreted protein that is involved in the regulation of cell growth and differentiation."
  ]
},
{
  "question": "Are long non coding RNAs spliced?",
  "goldstandard_documents": [
    "Long non coding RNAs are not spliced."
  ]
},
{
  "question": "Is RANKL secreted from the cells?",
  "goldstandard_documents": [
    "RANKL is a secreted protein that is involved in the regulation of bone metabolism."
  ]
},
{
  "question": "Which miRNAs could be used as potential biomarkers for epithelial ovarian cancer?",
  "goldstandard_documents": [
    "miR-21, miR-200c, and miR-141 are potential biomarkers for epithelial ovarian cancer."
  ]
},
{
  "question": "Which acetylcholinesterase inhibitors are used for treatment of myasthenia gravis?",
  "goldstandard_documents": [
    "Pyridostigmine and neostigmine are used for treatment of myasthenia gravis."
  ]
},
{
  "question": "Has Denosumab (Prolia) been approved by FDA?",
  "goldstandard_documents": [
    "Denosumab (Prolia) has been approved by the FDA for the treatment of osteoporosis."
  ]
},
{
  "question": "List the human genes encoding for the dishevelled proteins?",
  "goldstandard_documents": [
    "Dishevelled 1, Dishevelled 2, and Dishevelled 3 are human genes encoding for the dishevelled proteins."
  ]
},
{
  "question": "Name synonym of Acrokeratosis paraneoplastica.",
  "goldstandard_documents": [
    "Acrokeratosis paraneoplastica is also known as Bazex syndrome."
  ]
}
```

LLM




```
{
  "question": "List the human genes encoding for the dishevelled proteins?",
  "goldstandard_documents": [
    "Dishevelled 1, Dishevelled 2, and Dishevelled 3 are human genes encoding for the dishevelled proteins."
  ]
},
{
  "question": "Name synonym of Acrokeratosis paraneoplastica.",
  "goldstandard_documents": [
    "Acrokeratosis paraneoplastica is also known as Bazex syndrome."
  ]
}
```


How to train a neural IR model?

❖ Conditions for a good synthetic data

- Query-document pair needs to be relevant.
- The query needs to be a good example, i.e. suitable for retrieval (stand alone questions).
- Example of bad queries:
 - Who is the author of this paper?
 - Who wrote this article



NIR Training – Optimization

❖ Pointwise

- Simple classification problem, we try to predict if a query-document pair is relevant or not according to the GS. Ex: Binary Cross Entropy Loss

❖ Pairwise

- Ordering problem, we try to predict where the query-document pair is more relevant than the negative. Ex: hinge loss, RankNet loss

❖ Listwise

- Optimizes the entire list, given the full query-document pairs we try to optimize the entire list such that the positive document have the highest scores. Ex: ListNet

NIR Training – Put everything together

- ❖ Data Loading and prepare the data loaders
 - Positive and negative examples
- ❖ Neural model instantiation
- ❖ Defining training objective
- ❖ Run training loop
 - Save the model checkpoints during training and register metrics

```
import torch
from torch import nn

# 1. Data Loading (text ids for query and document)
train_loader = DataLoader(
    dataset, # [(query_ids, doc_ids, labels), ...]
    batch_size=32,
    shuffle=True
)

# 2. IR Neural Model
class IRModel(nn.Module):
    ...

# 3. Training objective
loss_function = nn.BCELoss()
model = IRModel()
optimizer = torch.optim.Adam(model.parameters())

# 4. Training Loop
optimizer = torch.optim.Adam(model.parameters())
for epoch in range(epochs):
    for query_ids, doc_ids, labels in train_loader:
        pred_scores = model(query_ids, doc_ids)
        loss = loss_function(pred_scores, labels)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

NIR Training – Put everything together

- ❖ Data Loading and prepare the data loaders
 - Positive and negative examples
- ❖ Neural model instantiation
- ❖ Defining training objective
- ❖ Run training loop
 - Save the model checkpoints during training and register metrics

```
import torch
from torch import nn

# 1. Data Loading (text ids for query and document)
train_loader = DataLoader(
    dataset, # [(query_ids, doc_ids, labels), ...]
    batch_size=32,
    shuffle=True
)


# 2. IR Neural Model
class IRModel(nn.Module):
    ...

# 3. Training objective
loss_function = nn.BCELoss()
model = IRModel()
optimizer = torch.optim.Adam(model.parameters())


# 4. Training Loop
trainer = Trainer(train_loader, model, loss_function)
trainer.train()
```



Literature Examples




Literature Examples



Shallow representation-based

Representation-Based



Average Word Embedding

- ❖ Simplest NIR model that creates dense representations by averaging the query and document terms embeddings.


$$\vec{q}^e = \frac{1}{|\vec{q}|} \sum_{i=1}^{|\vec{q}|} \frac{\vec{u}_i}{\|\vec{u}_i\|_2},$$

$$\vec{d}^e = \frac{1}{|\vec{d}|} \sum_{i=1}^{|\vec{d}|} \frac{\vec{v}_i}{\|\vec{v}_i\|_2}.$$

- ❖ Adding a tf-idf weighting scheme

$$\vec{q}^e = \frac{\sum_{i=1}^{|\vec{q}|} \frac{\vec{u}_i}{\|\vec{u}_i\|} \times tf(u_i, \vec{q}) \times idf(u_i)}{\sum_{i=1}^{|\vec{q}|} tf(u_i, \vec{q}) \times idf(u_i)},$$

$$\vec{d}^e = \frac{\sum_{i=1}^{|\vec{d}|} \frac{\vec{v}_i}{\|\vec{v}_i\|} \times tf(v_i, \vec{d}) \times idf(v_i)}{\sum_{i=1}^{|\vec{d}|} tf(v_i, \vec{d}) \times idf(v_i)}.$$



Average Word Embedding

- ❖ Simplest NIR model that creates dense representations by averaging the query and document terms embeddings.


$$\vec{q}^e = \frac{1}{|\vec{q}|} \sum_{i=1}^{|\vec{q}|} \frac{\vec{u}_i}{\|\vec{u}_i\|_2},$$

$$\vec{d}^e = \frac{1}{|\vec{d}|} \sum_{i=1}^{|\vec{d}|} \frac{\vec{v}_i}{\|\vec{v}_i\|_2}.$$

- ❖ Adding a tf-idf weighting scheme

$$\vec{q}^e = \frac{\sum_{i=1}^{|\vec{q}|} \frac{\vec{u}_i}{\|\vec{u}_i\|} \times tf(u_i, \vec{q}) \times idf(u_i)}{\sum_{i=1}^{|\vec{q}|} tf(u_i, \vec{q}) \times idf(u_i)},$$

$$\vec{d}^e = \frac{\sum_{i=1}^{|\vec{d}|} \frac{\vec{v}_i}{\|\vec{v}_i\|} \times tf(v_i, \vec{d}) \times idf(v_i)}{\sum_{i=1}^{|\vec{d}|} tf(v_i, \vec{d}) \times idf(v_i)}.$$



It utilizes pretrained word embeddings. Optionally you can finetune these embeddings

DSSM (Deep Semantic Similarity Model)



Posterior probability
computed by softmax

Relevance measured
by cosine similarity

Semantic feature

y

$\{W_4, b_4\}$

l_3

Multi-layer non-
linear projection

$\{W_3, b_3\}$

l_2

Word Hashing


$\{W_2, b_2\}$

l_1

Term Vector


x

W_1




DSSM (Deep Semantic Similarity Model)

- ❖ Other variants that explored embeddings with:
 - Convolutional layers and LSTM layers
- ❖ Overall poor performance (vs BM25)





Shallow interaction-based





PACRR (2017)

- ❖ Mainly uses convolution kernels to extract matching patterns over a similarity matrix
- ❖ Adds a late IDF feature before the final ranking score computation

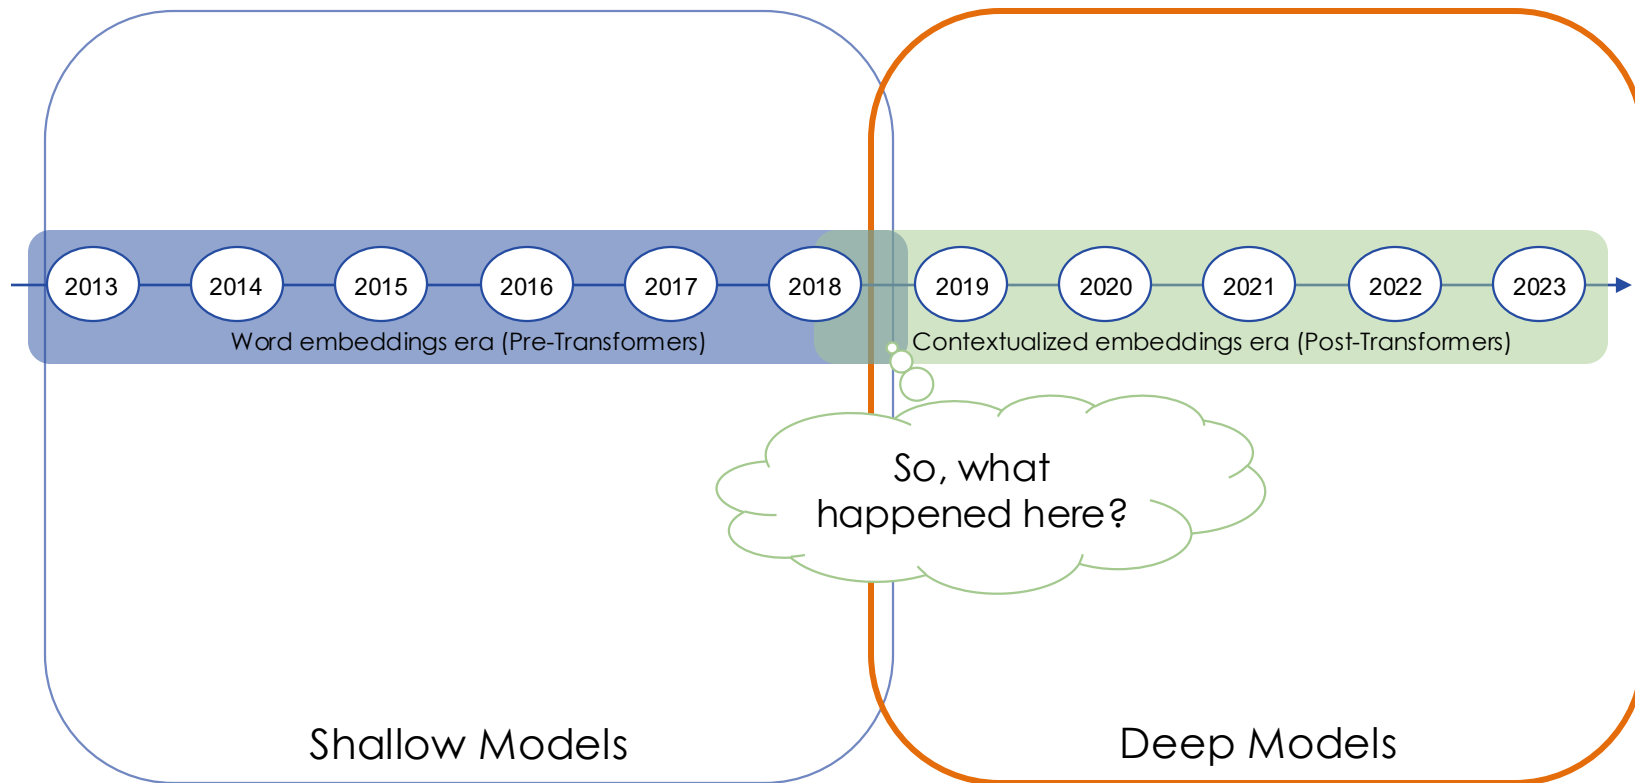


DeepRank (2017)


- ❖ Tries to mimic the way humans search for information.
- ❖ Weights the importance of each query term, since different query terms carry different importance for the information need.



Moving to the Deep Models



Moving to the Deep Models



Backbone (magic) behind LLMs (chatGPT)

Contextualize Transformer Models

- ❖ They revolutionize language modelling by efficiently processing and learning from billions of documents and web pages.
- ❖ They enable the creation of highly **contextualized** embeddings, capturing nuanced meanings and relationships in large-scale textual data.

Contextualize Transformer Models

Static word embeddings

I am opening a
bank account

I am sited on
the river bank

$\vec{w}(\text{bank})$

Contextualized word embeddings

I am opening a
bank account

I am sited on
the river bank

$\vec{w}(\text{bank})$

$\vec{w}(\text{bank})$

From previous lecture, do you remember this?

- ❖ Turns out creating contextualized word embeddings is hard, since it requires processing thousands of billions of tokens! Something that previous RNN based solutions were not mature (efficient) enough.

Contextualize Transformer Models

- ❖ Transformer models are larger models that highly scalable and can easily being training in parallel taking full advantage of accelerators like GPUs.
- ❖ The “downside” is that they are larger models that requires a lot of computational resources:
 - Some research groups use, for example, 2048, 1024 GPUs, 512 TPUs to produce their experiments.

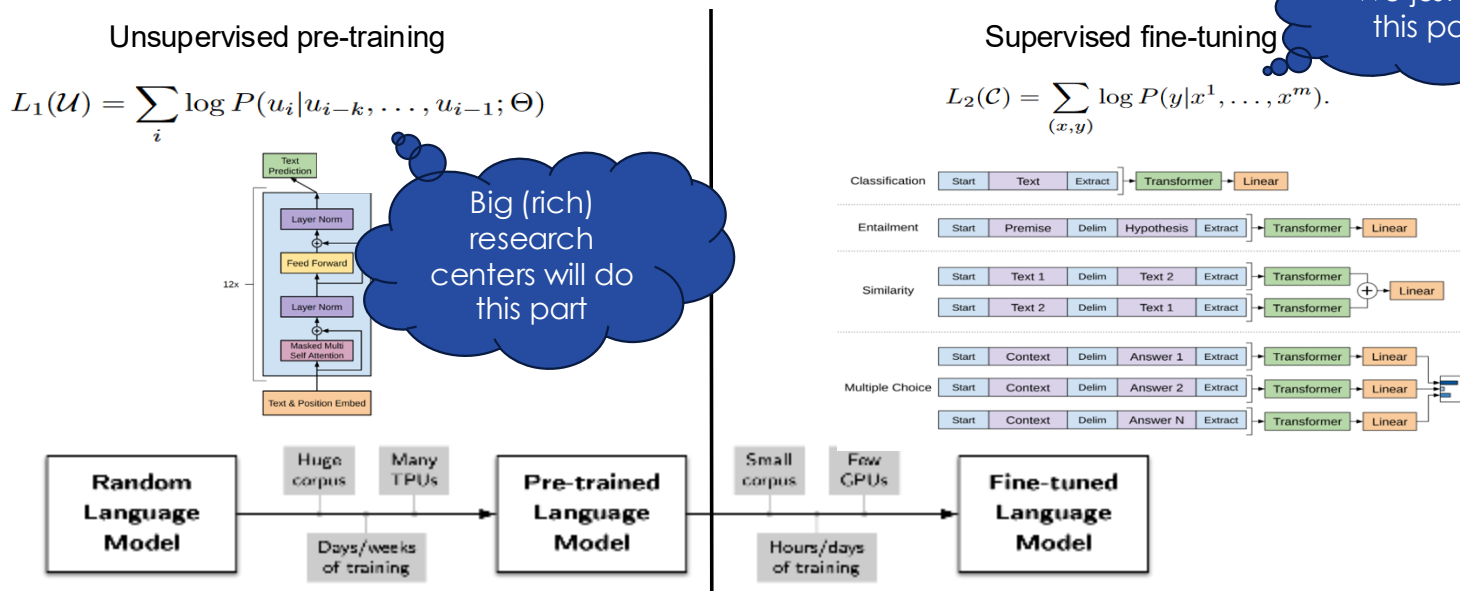


FYI: For today standard of LLMs this is nothing



Contextualize Transformer Models

- ❖ Thankfully, we can reuse these large models and just make small adjustments for our tasks.



Transformers Architectures

❖ 3 main architectures:

– Decoder only (LM)

- GPT
- Llama

– Encoder-Decoder

- T5

– Encoder only (MLM)

- BERT

Incapable of generating text,
but good for creating
contextualized representation
and for classification tasks.

Good for generating text

(2015, Dai and Le) TL in NLP

(2017, Vaswani et al.) Intro of the TB

(2018, Howard and Ruder) ULMFIT: TL in NLP

(2018, Radford et al.) GPT-1: TL + TB

(2018, Devlin et al.) BERT

(2019, Dai et al.) Transformer-XL

(2019, Lample and Conneau.) XLM

(2019, Radford et al.) GPT-2

(2019, Yang et al.) XLNET

(2019, Liu et al.) RoBERTa

(2019, Raffel et al.) T5

(2019, Sanh et al.) distilBERT

(2019, Lan et al.) ALBERT

(2019, Sholehybi et al.) Megatron-LM

(2019, Keskar et al.) CTRL

(2020, Beltagy et al.) Longformer

(2020, Clark et al.) ELECTRA

(2020, Rosset et al.) Turing-NLG

(2020, Kitaev et al.) Reformer

(2020, Brown et al.) GPT-3

(2020, Wang et al.) Linformer

June 2020


I personally
gave up on
keeping track
around here

Transformer in IR

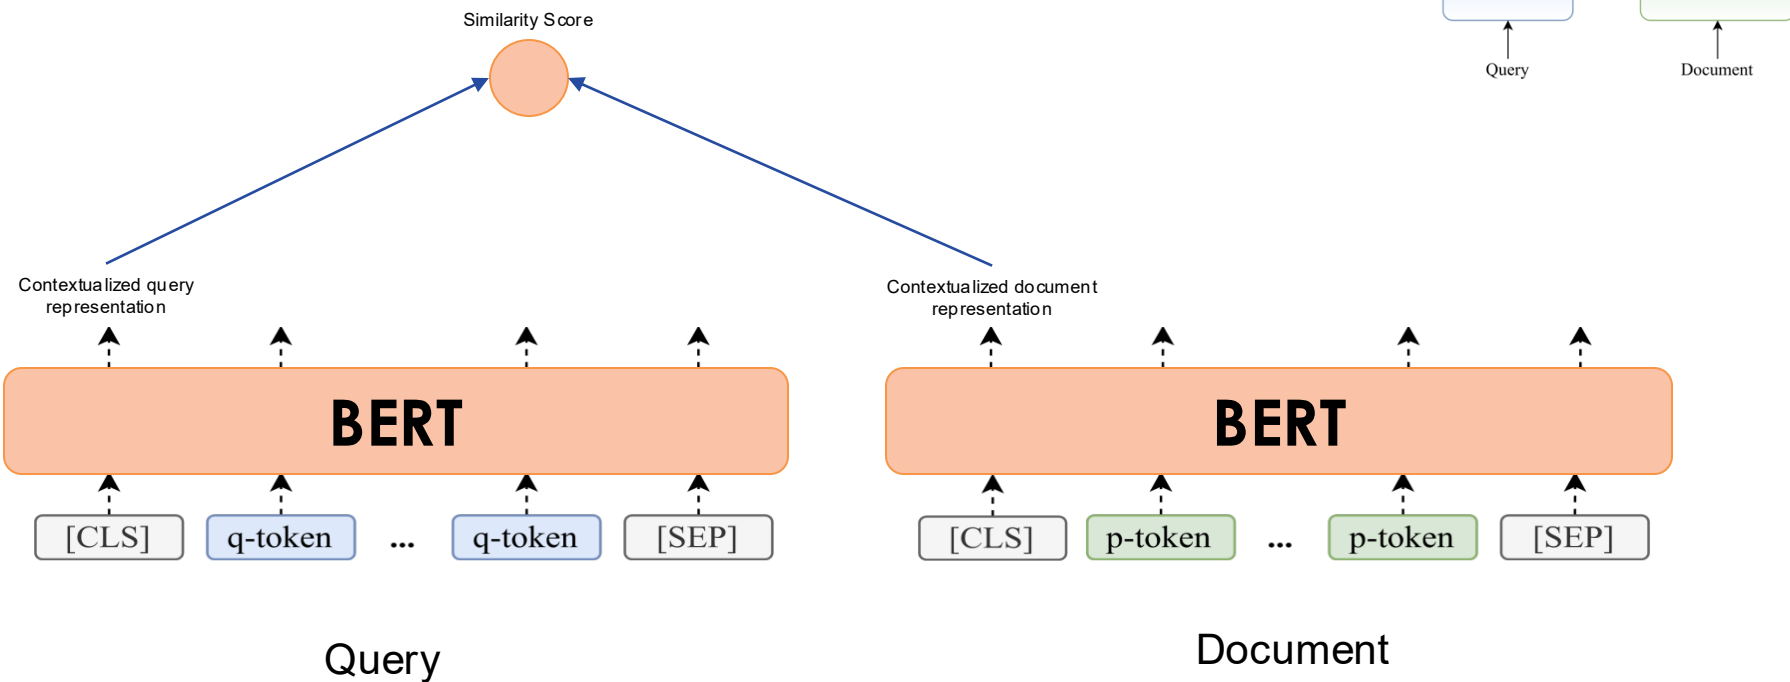
- ❖ In IR the encoder type is the mostly widely adopted transformer model to build neural information retrieval models.
- ❖ Main idea:
 - For representation-based, lets leverage the more powerful contextualized embeddings to build query and document representation
 - For interaction-based, lets frame the ranking problem as a classification problem, e.g., is my query relevant to my document (yes or no)

Deep representation-based


Representation-Based




DPR (Dense Passage Retrieval)




Deep interaction-based



Interaction-based models using Transformers




Mono and duo BERT



CEDR (Contextualized Embeddings for Document Ranking)

- ❖ Uses BERT to obtain the contextualized embeddings but then applies the previous neural IR models:
 - PACRR
 - KNRM
 - DRMM
 - DeepRank



Resources

- ❖ For transformer-based models having GPU necessary, for that:
 - Use colab from google the free tier offers GPU
 - Use Kaggle, the free tier offers 30h per week (more than enough)
- ❖ Pytorch: https://colab.research.google.com/drive/1KX9HUd--QupV_GdROvm3XA28ZyFCuscF?usp=sharing
- ❖ Embeddings: https://colab.research.google.com/drive/1OtLgsQD6wPNUcrRixOKAo_eXJL7rD2Px?usp=sharing